

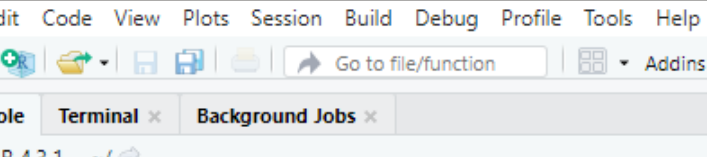
(LC3.1) What's another way using the "not" operator! to filter only the rows that are not going to Burlington, VT nor Seattle, WA in the flights data frame? Test this out using the code above.

```
not_BTV_SEA <- flights %>%
  filter(!(dest == "BTV" | dest == "SEA"))
not_BTV_SEA <- flights %>%
  filter(!dest == "BTV" & !dest == "SEA")
not_BTV_SEA <- flights %>%
  filter(dest != "BTV" & dest != "SEA")
```

(LC3.2) Say a doctor is studying the effect of smoking on lung cancer for a large number of patients who have records measured at five-year intervals. She notices that a large number of patients have missing data points because the patient has died, so she chooses to ignore these patients in her analysis. What is wrong with this doctor's approach?

Lung cancer may have killed the missing patients! So ignoring them might severely skew our results! It is critical to consider the implications of omitting missing data for our study! Considering the following:
Is there a reason why some values are missing? If this is the case, our findings may be skewed!
If there isn't, it may be acceptable to "sweep missing values under the rug."

(LC3.3) Modify the above summarize function to create summary_temp to also use the n() summary function: summarize(count = n()). What does the returned value correspond to?



The screenshot shows the RStudio application window. The top menu bar includes File, Edit, Code, View, Plots, Session, Build, Debug, Profile, Tools, and Help. Below the menu is a toolbar with icons for adding files, saving, and navigating. The main pane is divided into three tabs: Console, Terminal, and Background Jobs. The Console tab is active, showing the R prompt and the following code and output:

```
> library(dplyr)
> library(ggplot2)
> library(nycflights13)
> weather %>%
+   summarize(count = n())
# A tibble: 1 x 1
  count
  <int>
1 26115
> |
```

```
weather %>%  
  summarize(count = n())
```

```
# A tibble: 1 × 1
  count
<int>
1 26115
```

(LC3.4) Why doesn't the following code work? Run the code line by line instead of all at once, and then look at the data. In other words, run `summary_temp <- weather %>% summarize(mean = mean(temp, na.rm = TRUE))` first.

```

R 4.3.1 ~ /
> summary_temp <- weather %>%
+ summarize(mean = mean(temp, na.rm = TRUE)) %>%
+ summarize(std_dev = sd(temp, na.rm = TRUE))
Error in `summarize()`:
! In argument: `std_dev = sd(temp, na.rm = TRUE)`.
Caused by error:
! object 'temp' not found
Run `rlang::last_trace()` to see where the error occurred.
> weather %>%
+ summarize(mean = mean(temp, na.rm = TRUE))
# A tibble: 1 x 1
  mean
  <dbl>
1  55.3
> |

```

```

weather %>%
  summarize(mean = mean(temp, na.rm = TRUE))

```

```

# A tibble: 1 x 1
  mean
  <dbl>
1 55.3

```

As the variable `temp` has been compressed to the value `mean` after the first `summarise()`. When we try to perform the second `summarise()`, it is unable to locate the variable `temp` on which to compute the standard deviation.

(LC3.5) Recall from Chapter 2 when we looked at plots of temperatures by months in NYC. What does the standard deviation column in the `summary_monthly_temp` data frame tell us about temperatures in New York City throughout the year?

RStudio

File Edit Code View Plots Session Build Debug Profile

Go to file/function

flights × weather × by_monthly_origin × coun

Filter

| | month | mean | std_dev |
|----|-------|----------|-----------|
| 1 | 1 | 35.63566 | 10.224635 |
| 2 | 2 | 34.27060 | 6.982378 |
| 3 | 3 | 39.88007 | 6.249278 |
| 4 | 4 | 51.74564 | 8.786168 |
| 5 | 5 | 61.79500 | 9.681644 |
| 6 | 6 | 72.18400 | 7.546371 |
| 7 | 7 | 80.06622 | 7.119699 |
| 8 | 8 | 74.46847 | 5.191615 |
| 9 | 9 | 67.37129 | 8.465902 |
| 10 | 10 | 60.07113 | 8.846035 |
| 11 | 11 | 44.99043 | 10.443805 |
| 12 | 12 | 38.44180 | 9.982432 |

(LC3.6) What code would be required to get the mean and standard deviation temperature for each day in 2013 for NYC?

```
summary_temp_by_day <- weather %>%
  group_by(year, month, day) %>%
  summarize(
    mean = mean(temp, na.rm = TRUE),
    std_dev = sd(temp, na.rm = TRUE)
  )
summary_temp_by_day
```

The screenshot shows the RStudio interface with the following code in the console:

```
> view(weather)
> library(dplyr)
> library(ggplot2)
> library(nycflights13)
> summary_temp_by_day <- weather %>%
+   group_by(year, month, day) %>%
+   summarize(
+     mean = mean(temp, na.rm = TRUE),
+     std_dev = sd(temp, na.rm = TRUE)
+   )
`summarize()` has grouped output by 'year', 'month'. You can override using the `.groups` argument.
> summary_temp_by_day
# A tibble: 364 x 5
# Groups:   year, month [12]
   year month   day mean std_dev
  <int> <int> <int> <dbl> <dbl>
1  2013     1     1  37.0    4.00
2  2013     1     2  28.7    3.45
3  2013     1     3  30.0    2.58
4  2013     1     4  34.9    2.45
5  2013     1     5  37.2    4.01
6  2013     1     6  40.1    4.40
7  2013     1     7  40.6    3.68
8  2013     1     8  40.1    5.77
9  2013     1     9  43.2    5.40
10 2013     1    10  43.8    2.95
# i 354 more rows
# i Use `print(n = ...)` to see more rows
> |
```

(LC3.7) Recreate by_monthly_origin, but instead of grouping via group_by(origin, month), group variables in a different order group_by(month, origin). What differs in the resulting dataset?

```
count_flights_by_airport <- flights %>%
  group_by(origin, carrier) %>%
  summarize(count = n())
```

| | month | origin | count |
|----|-------|--------|-------|
| 1 | 1 | EW R | 9893 |
| 2 | 1 | JFK | 9161 |
| 3 | 1 | LGA | 7950 |
| 4 | 2 | EW R | 9107 |
| 5 | 2 | JFK | 8421 |
| 6 | 2 | LGA | 7423 |
| 7 | 3 | EW R | 10420 |
| 8 | 3 | JFK | 9697 |
| 9 | 3 | LGA | 8717 |
| 10 | 4 | EW R | 10531 |
| 11 | 4 | JFK | 9218 |
| 12 | 4 | LGA | 8581 |
| 13 | 5 | EW R | 10592 |
| 14 | 5 | JFK | 9397 |
| 15 | 5 | LGA | 8807 |
| 16 | 6 | EW R | 10175 |
| 17 | 6 | JFK | 9472 |
| 18 | 6 | LGA | 8596 |
| 19 | 7 | EW R | 10475 |
| 20 | 7 | JFK | 10023 |
| 21 | 7 | LGA | 8927 |
| 22 | 8 | EW R | 10359 |
| 23 | 8 | JFK | 9983 |
| 24 | 8 | LGA | 8985 |
| 25 | 9 | EW R | 9550 |
| 26 | 9 | JFK | 8886 |

```

R 4.3.1 ~ /
> by_monthly_origin <- flights %>%
+   group_by(month, origin) %>%
+   summarize(count = n())
`summarise()` has grouped output by 'month'. You can override using the `.groups` argument.
> view(by_monthly_origin)
> |

```

The month column is now first in `by_monthly_origin`, and the rows are ordered by month rather than origin. When we use the `View()` method to compare the values of count in `by_origin_monthly` and `by_monthly_origin`, we'll see that the data are the same, only displayed in a different order.

(LC3.8) How could we identify how many flights left each of the three airports for each carrier?

Using the `n()` method, which counts rows, we may summarise the count from each airport.

```

count_flights_by_airport <- flights %>%
  group_by(origin, carrier) %>%
  summarize(count = n())

```

| | origin | carrier | count |
|----|--------|---------|-------|
| 1 | EWR | 9E | 1268 |
| 2 | EWR | AA | 3487 |
| 3 | EWR | AS | 714 |
| 4 | EWR | B6 | 6557 |
| 5 | EWR | DL | 4342 |
| 6 | EWR | EV | 43939 |
| 7 | EWR | MQ | 2276 |
| 8 | EWR | OO | 6 |
| 9 | EWR | UA | 46087 |
| 10 | EWR | US | 4405 |
| 11 | EWR | VX | 1566 |
| 12 | EWR | WN | 6188 |
| 13 | JFK | 9E | 14651 |
| 14 | JFK | AA | 13783 |
| 15 | JFK | B6 | 42076 |
| 16 | JFK | DL | 20701 |
| 17 | JFK | EV | 1408 |
| 18 | JFK | HA | 342 |
| 19 | JFK | MQ | 7193 |
| 20 | IFK | LJA | 4534 |

(LC3.9) How does the filter operation differ from a group by followed by a summarize?

The filter extracts rows from the original dataset while leaving the others alone, although the group by%>% summarise generates new values by computing summaries of numerical variables.

(LC3.10) What do positive values of the gain variable in flights correspond to? What about negative values? And what about a zero value?

Let's assume that a flight was delayed by 30 minutes, `dep_delay = 30`.
came after that 20 minutes late, or `arr_delay = 20`.

As a result, `gain = dep_delay - arr_delay = 30 - 20 = 10` is positive, indicating that it "made up/gained time in the air."

If both the departure and arrival times were 0, no extra time was added. We see that the increase is typically close to zero minutes.

(LC3.11) Could we create the `dep_delay` and `arr_delay` columns by simply subtracting `dep_time` from `sched_dep_time` and similarly for arrivals? Try the code out and explain any differences between the result and what actually appears in flights

No, because times cannot be directly arithmetic. There are 4 minutes between 12:03 and 11:59, yet 12:03 and 11:59 equal 44.

(LC3.12) What can we say about the distribution of gain? Describe it in a few sentences using the plot and the `gain_summary` data frame values.

The gain is often between -50 and 50 minutes and slightly over zero (the median is 7, implying gain is above 0 at least 50% of the time). Although, there are some extreme instances!

(LC3.13) Looking at Figure 3.7, when joining flights and weather (or, in other words, matching the hourly weather values with each flight), why do we need to join by all of year, month, day, hour, and origin, and not just hour?

Hour is only a number between 0 and 23, thus we need to know the year, month, day, and airport in order to identify a certain hour.

(LC3.14) What surprises you about the top 10 destinations from NYC in 2013?

The large number of flights to Boston surprised me.

(LC3.15) What are some advantages of data in normal forms? What are some disadvantages?

We can simply join datasets with other datasets when they are in normal form! We could, for instance, combine flight and aircraft data.

(LC3.16) What are some ways to select all three of the `dest`, `air_time`, and `distance` variables from flights? Give the code showing how to do this in at least three different ways.

```
flights %>%  
  select(dest, air_time, distance)
```

```
flights %>%  
+   select(dest:distance)
```

```
flights %>%  
  select(  
    -year, -month, -day, -dep_time, -sched_dep_time, -dep_delay, -arr_time,  
    -sched_arr_time, -arr_delay, -carrier, -flight, -tailnum, -origin,  
    -hour, -minute, -time_hour  
  )
```

Showing 1 to 11 of 35 entries, 3 total columns

```
Console Terminal Background Jobs
R 4.3.1 ~ /
> flights %>%
+   select(dest, air_time, distance)
# A tibble: 336,776 x 3
  dest   air_time distance
<chr>   <dbl>   <dbl>
1 IAH     227     1400
2 IAH     227     1416
3 MIA     160     1089
4 BQN     183     1576
5 ATL     116     762
6 ORD     150     719
7 FLL     158     1065
8 IAD      53      229
9 MCO     140     944
10 ORD     138     733
# i 336,766 more rows
# i Use `print(n = ...)` to see more rows
>
```

```
> flights %>%
+   select(dest:distance)
# A tibble: 336,776 x 3
  dest   air_time distance
<chr>   <dbl>   <dbl>
1 IAH     227     1400
2 IAH     227     1416
3 MIA     160     1089
4 BQN     183     1576
5 ATL     116     762
6 ORD     150     719
7 FLL     158     1065
8 IAD      53      229
9 MCO     140     944
10 ORD     138     733
# i 336,766 more rows
# i Use `print(n = ...)` to see more rows
> |
```

```
> flights %>%
+   select(
+     -year, -month, -day, -dep_time, -sched_dep_time, -dep_delay, -arr_time,
+     -sched_arr_time, -arr_delay, -carrier, -flight, -tailnum, -origin,
+     -hour, -minute, -time_hour
+   )
# A tibble: 336,776 x 3
  dest   air_time distance
<chr>   <dbl>   <dbl>
1 IAH     227     1400
2 IAH     227     1416
3 MIA     160     1089
4 BQN     183     1576
5 ATL     116     762
6 ORD     150     719
7 FLL     158     1065
8 IAD      53      229
9 MCO     140     944
10 ORD     138     733
# i 336,766 more rows
# i Use `print(n = ...)` to see more rows
> |
```

(LC3.17) How could one use `starts_with`, `ends_with`, and `contains` to select columns from the `flights` data frame? Provide three different examples in total: one for `starts_with`, one for `ends_with`, and one for `contains`

```
flights %>%
  select(starts_with("d"))
```

```
flights %>%
  select(ends_with("delay"))
```

```
flights %>%
  select(contains("dep"))
```

```
> flights %>%
+   select(starts_with("d"))
# A tibble: 336,776 x 5
  day dep_time dep_delay dest distance
<int> <int> <dbl> <chr> <dbl>
1     1     517         2 IAH     1400
2     1     533         4 IAH     1416
3     1     542         2 MIA     1089
4     1     544        -1 BQN     1576
5     1     554        -6 ATL      762
6     1     554        -4 ORD      719
7     1     555        -5 FLL     1065
8     1     557        -3 IAD      229
9     1     557        -3 MCO      944
10    1     558        -2 ORD      733
# i 336,766 more rows
# i Use `print(n = ...)` to see more rows
> |
```

```
> flights %>%
+   select(ends_with("delay"))
# A tibble: 336,776 x 2
  dep_delay arr_delay
<dbl> <dbl>
1         2         11
2         4         20
3         2         33
4        -1        -18
5        -6        -25
6        -4         12
7        -5         19
8        -3        -14
9        -3         -8
10       -2          8
# i 336,766 more rows
# i Use `print(n = ...)` to see more rows
> |
```

```
> flights %>%
+   select(contains("dep"))
# A tibble: 336,776 x 3
  dep_time sched_dep_time dep_delay
<int> <int> <dbl>
1     517         515         2
2     533         529         4
3     542         540         2
4     544         545        -1
5     554         600        -6
6     554         558        -4
7     555         600        -5
8     557         600        -3
9     557         600        -3
10    558         600        -2
# i 336,766 more rows
# i Use `print(n = ...)` to see more rows
> |
```

(LC3.18) Why might we want to use the select() function on a data frame?

To reduce the data frame's size and make it easier to examine. use View() as an illustration.

(LC3.19) Create a new data frame that shows the top 5 airports with the largest arrival delays from NYC in 2013.

```
top_five <- flights %>%
  group_by(dest) %>%
  summarize(avg_delay = mean(arr_delay, na.rm = TRUE)) %>%
  arrange(desc(avg_delay)) %>%
  top_n(n = 5)
top_five
```



```

Console Terminal Background Jobs
R 4.3.1 . ~/
> top_five <- flights %>%
+   group_by(dest) %>%
+   summarize(avg_delay = mean(arr_delay, na.rm = TRUE)) %>%
+   arrange(desc(avg_delay)) %>%
+   top_n(n = 5)
Selecting by avg_delay
> top_five
# A tibble: 5 x 2
  dest avg_delay
<chr>   <dbl>
1 CAE     41.8
2 TUL     33.7
3 OKC     30.6
4 JAC     28.1
5 TYS     24.1
>

```

(LC3.20) Using the datasets included in the nycflights13 package, compute the available seat miles for each airline sorted in descending order. After completing all the necessary data wrangling steps, the resulting data frame should have 16 rows (one for each airline) and 2 columns (airline name and available seat miles). Here are some hints:

```

flights %>%
  inner_join(planes, by = "tailnum") %>%
  select(carrier, seats, distance) %>%
  mutate(ASM = seats * distance) %>%
  group_by(carrier) %>%
  summarize(ASM = sum(ASM, na.rm = TRUE)) %>%
  arrange(desc(ASM))

```

```

Showing 1 to 11 of 35 entries, 3 total columns
Console Terminal Background Jobs
R 4.3.1 . ~/
> flights %>%
+   inner_join(planes, by = "tailnum") %>%
+   select(carrier, seats, distance) %>%
+   mutate(ASM = seats * distance) %>%
+   group_by(carrier) %>%
+   summarize(ASM = sum(ASM, na.rm = TRUE)) %>%
+   arrange(desc(ASM))
# A tibble: 16 x 2
  carrier      ASM
<chr>      <dbl>
1 UA    15516377526
2 DL    10532885801
3 B6     9618222135
4 AA     3677292231
5 US     2533505829
6 VX     2296680778
7 EV     1817236275
8 WN     1718116857
9 9E      776970310
10 HA     642478122
11 AS     314104736
12 FL     219628520
13 F9     184832280
14 YV     20163632
15 MQ      7162420
16 OO     1299835
> |

```